

Credit Risk Modelling For Dummies: But With Fewer Dummies

S. J. Curran

Overview

Standard credit risk analysis utilises scorecards which are built using only datasets with categorical variables. This requires the continuous numerical features to be fine-classed (grouped into discrete sets) and converted to dummy variables. Although the point of the scorecard is to present the model in a simple way, this practice requires much convoluted pre-processing of the data, which greatly bloats the size of the dataset and makes it more susceptible to containing errors. Most importantly though, I find that, by retaining the numerical features, the predictive power of the data is great improved, with a Gini coefficient of 0.95 (cf. 0.40 with fine-classing) and Kolmogorov-Smirnov statistic of $KS = 0.85$ (cf. 0.30).

Although all of the processing was done in the python scripts mentioned in the report, a simplified run-through is included as the Jupyter notebooks [PP.ipynb](#) and [ML.ipynb](#)

Contents

1	Introduction	1
2	Pre-processing	1
2.1	Initial pre-processing	1
2.1.1	The data	1
2.1.2	Times and dates	2
2.1.3	Further trimming	4
2.1.4	Imputing the remaining missing values	5
2.1.5	Formatting the target variable	6
2.2	Coarse classing of categorical values	11
2.2.1	Grade	11
2.2.2	Home ownership	13
2.2.3	Verification status	13

2.2.4	Payment plan	14
2.2.5	Purpose	14
2.2.6	Address state	15
2.2.7	Initial list status	17
3	Machine Learning	18
3.1	Initial models and tests	18
3.2	Model optimisation	19
3.3	Feature importance	20
4	Results	21
4.1	ROC curve	21
4.2	Gini coefficient	22
4.3	Kolmogorov-Smirnov score	22
4.4	Using the Udemmy Target Classification	23
4.5	IVs of the Most Important Features	24

1 Introduction

Credit risk addresses the potential loss a lender or investor may face if a borrower or debtor fails to repay a loan or meet their financial obligations. It is a critical consideration for financial institutions, as a high level of credit risk can lead to significant financial losses and even bankruptcy. Accurately assessing credit risk is essential for lenders to make informed decisions about whether to extend credit and the terms of the loan.

Coarse and fine classing are methods used to assess credit risk. *Coarse classing* involves grouping borrowers or debtors into broad categories based on their creditworthiness, such as high, medium, or low risk. *Fine classing*, on the other hand, involves a more detailed analysis of credit risk, using a more nuanced approach to identify different levels of creditworthiness within each broad category.

Classing is used as credit regulators require statistical models for estimating credit risk to be presented in a simple way (the scorecard). This is based upon the coefficients in the *probability of default* model and must only contain only dummy variables as predictive features. However, fine classing (which generates typically 20 – 50 features) can lead quickly to a very bloated dataset. This requires more processing memory and is effectively an unnecessary degradation of any continuous (non-discrete data), while making the pre-processing of data more complex and thus more prone to errors.

Here I demonstrate that not fine classing can deliver a significantly greater predictive power¹, while simplifying the pre-processing of the data.

2 Pre-processing

2.1 Initial pre-processing

2.1.1 The data

The dataset `loan_data_200_2014.csv` (available from [Kaggle](#)) contains all available data for more than 800 000 consumer loans issued from 2007 to 2015 by LendingClub, a large U.S. peer to peer lending company.

Running `pre-process_1.py`, it is seen that there are 466285 rows x 75 columns. Also, it is evident that there are a lot of missing values (NaN) and a lot of feature columns which are completely empty, e.g.

```
54 annual_inc_joint          0 non-null    float64
55 dti_joint                 0 non-null    float64
56 verification_status_joint 0 non-null    float64
```

To start with, I drop the features with > 200 000 missing values (a large fraction of the data). These are

```
desc has 340302 missing values
```

¹Compared to Udemy's [Credit Risk Modelling in Python](#) and [Credit Risk Modelling \(Part II\)](#).

```

mths_since_last_delinq has 250351 missing values
mths_since_last_record has 403647 missing values
next_pymnt_d has 227214 missing values
mths_since_last_major_derog has 367311 missing values
annual_inc_joint has 466285 missing values
dti_joint has 466285 missing values
verification_status_joint has 466285 missing values
open_acc_6m has 466285 missing values
open_il_6m has 466285 missing values
open_il_12m has 466285 missing values
open_il_24m has 466285 missing values
mths_since_rcnt_il has 466285 missing values
total_bal_il has 466285 missing values
il_util has 466285 missing values
open_rv_12m has 466285 missing values
open_rv_24m has 466285 missing values
max_bal_bc has 466285 missing values
all_util has 466285 missing values
inq_fi has 466285 missing values
total_cu_tl has 466285 missing values
inq_last_12m has 466285 missing values

```

The extraneous fields – Unnamed: 0 and member_id are also removed. At this stage, id (a unique ID for the loan) is retained as this will be useful in recombining the data if this is split (e.g. Sect. 2.1.4). After removal there are 50 feature columns remaining.

2.1.2 Times and dates

Following the [Udemy](#) analysis, the next step is to covert the *employment length* (emp_length) and *term of the loan* (term) to numeric values. That is

```

['10+ years' '< 1 year' '1 year' '3 years' '8 years' '9 years' '4 years'
 '5 years' '6 years' '2 years' '7 years' nan]

[' 36 months' ' 60 months']

```

which become

```

['10' '0' '1' '3' '8' '9' '4' '5' '6' '2' '7' nan]
[36 60]

```

In order to limit empty space between the columns when viewing the data-frame, I have also shortened the feature name emp_length to EM_L.

The *earliest credit line* (earliest_cr_line_date) field also has to be converted into a useful format, and using

```
pd.to_datetime(df['earliest_cr_line'],
               format = '%b-%y')
```

gives the second column to the right from the original first. I also shorten earliest_cr_line_date to ECL.

0	Jan-85	0	1985-01-01
1	Apr-99	1	1999-04-01
2	Nov-01	2	2001-11-01
3	Feb-96	3	1996-02-01
4	Jan-96	4	1996-01-01
...
466280	Apr-03	466280	2003-04-01
466281	Jun-97	466281	1997-06-01
466282	Dec-01	466282	2001-12-01
466283	Feb-03	466283	2003-02-01
466284	Feb-00	466284	2000-02-01

However, looking at the converted dates, the latest is over 40 years in the future! While this suggests super predictive power it is probably inadvisable to rely on this. As the UdeMy course suggests, this is due to the built-in time scale starting from 1970 (it actually looks like the start of 1969).

```
count      466256
unique      664
top      2000-10-01 00:00:00
freq      3674
first      1969-01-01 00:00:00
last      2068-12-01 00:00:00
Name: ECL, dtype: object
```

Looking at the latest earliest_cr_line_date (ECL) before today

```
from datetime import date
today = np.datetime64(date.today()); print(today)
temp = df[df['ECL'] < today]; print(temp['ECL'].describe())
```

gives 2011-11-01 as the latest ECL date.

	count	466256.000000
	mean	166.492177
	std	93.976747
	min	-685.024333
	max	513.981807

The earliest credit line is then converted to months since this date (UdeMy uses December 2017), giving the output to the right. The error in the dates seen above is apparent from the minimum number of months since the earliest credit line.

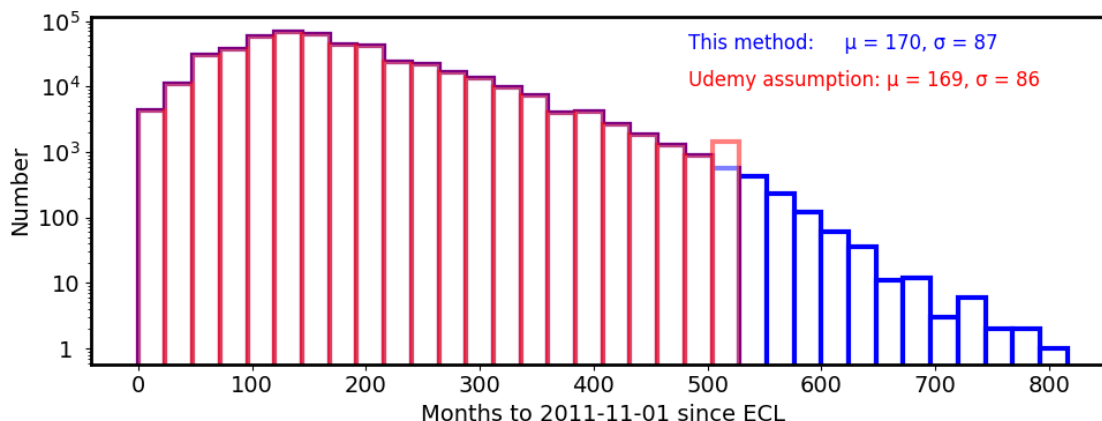
The [UdeMy](#) course deals with this by substituting the negative values for the maximum observed positive difference, with the justification “*Even if we don’t calculate the exact number of months that have passed since the earliest credit line was issued for those issued in the 60s, we put a very large value and we still get pretty close to the real picture.*”

However, it is arguable whether these values are *very large*, as well as the potential of unnecessarily losing crucial information. By splitting the date string and subtracting 100 from the year for which the number of months are negative, we can obtain the correct number of months since the earliest credit line.

count	466256.000000	count	466256.000000
mean	169.395094	mean	169.500887
std	86.461715	std	86.930703

min	0.000000	min	0.000000
max	513.981807	max	814.012608

From the mean values and the plot below, it is seen that the Udeemy assumption should have minimum impact, although this could be an issue if working with a significant fraction of older (pre-1970) loans. So while the assumption is easier to implement, it may not always be warranted.



This procedure is repeated to get the months since the loan issue date (`issue_d`). This spans 2007-06-01 to 2014-12-01 and so no conversion of pre-1970 dates is required. This is also the case for `last_pymnt_d` (2007-12-01 to 2016-01-01) and `last_credit_pull_d` (2007-12-01 to 2016-01-01), which are also converted.

```
count      466256
unique           91
top    2014-10-01 00:00:00
freq      38782
first    2007-06-01 00:00:00
last     2014-12-01 00:00:00
Name: i_date, dtype: object
```

2.1.3 Further trimming

Before dealing with the categorical features, the data which appear to be of limited value are further trimmed:

- `url` – as these are unique to each customer²
- `zip_code` – which are only partial (e.g. 860xx) and we will use the *address state* (`addr_state`)
- `sub_grade` – as credit worthiness grade (`grade`) will be used, and perhaps coarse classed (Sect. 2.2.1) so that higher resolution values will not be needed
- `application_type` – which only has a single value (namely INDIVIDUAL)

Revisiting the missing values

²Of the form [https://www.lendingclub.com/browse/loanDetail.action?loan_id= ...](https://www.lendingclub.com/browse/loanDetail.action?loan_id=...)

```

emp_title has 27541 missing values
title has 20 missing values
revol_util has 305 missing values
collections_12_mths_ex_med has 115 missing values
tot_coll_amt has 70130 missing values
tot_cur_bal has 70130 missing values
total_rev_hi_lim has 70130 missing values
EM_L has 20983 missing values

```

the employment length (EM_L) could be important, as well as the borrower's job title (emp_title), so I initially removed only those with > 30 000 missing values. Of the remaining missing values:

- emp_title has 205 271 unique values, but these are categorical, which would require too many dummy variables, so this is *removed*.
- title has 63015 unique values which are also categorical, so this is *removed*.
- revol_util has 1270 unique values which are numerical, so this will be *imputed*.
- collections_12_mths_ex_med has 10 unique values which are numerical, so this will be *imputed*.
- EM_L has 12 unique values which are numerical, so this will be *imputed*.

Which now leaves 41 features in the data.

2.1.4 Imputing the remaining missing values

There are various methods available to impute missing data, with the most common being *simple imputation*, where the missing feature is replaced with either a constant, the mean, the median or the most frequently occurring value. Since this is basically just using an assumed value, I opted for *multivariate (multiple) imputation*, where machine learning is used to estimate the missing values from the other features.

Bearing in mind the caveats that the employment length EM_L contains both lower (10+ years) and upper limits (< 1 year), I summarise the effects of the imputation below.

```

----- BEFORE IMPUTATION -----
For revol_util, n = 465534, mean = 56.181 +/- 23.730
For collections_12_mths_ex_med, n = 465724, mean = 0.009 +/- 0.109
For EM_L, n = 444856, mean = 5.994 +/- 3.627

----- AFTER IMPUTATION -----
For revol_util, n = 465839, mean = 56.186 +/- 23.725
For collections_12_mths_ex_med, n = 465839, mean = 0.009 +/- 0.109
For EM_L, n = 465839, mean = 6.019 +/- 3.555

```

2.1.5 Formatting the target variable

The target variable is the `loan_status`, which has the unique values

```
'Fully Paid'           'Late (31-120 days)'
'Charged Off'          'In Grace Period'
'Late (16-30 days)'    'Does not meet the credit policy. Status:Fully Paid'
'Current'              'Does not meet the credit policy. Status:Charged Off'
'Default'
```

of which, e.g. Fully Paid, Current and In Grace Period are considered good loans and Charged Off and Default as bad.

In order to classify these, I calculate the *debt-to-income ratio* (DTIR) for each category.³ To get the debt, I subtracted the *principal received to date* (`total_rec_prncp`) from the *loan amount* (`loan_amnt`). In increasing debt-to-income ratio, the results are

	<code>loan_status</code>	number	mean DTIR	standard error
0	Fully Paid	184724	0.0008	0.0
7	Does not meet the credit policy. Status:Fully ...	1961	0.0027	0.0004
8	Does not meet the credit policy. Status:Charge...	746	0.1194	0.0036
2	Current	224207	0.1283	0.0002
5	In Grace Period	3146	0.1399	0.0016
6	Late (16-30 days)	1218	0.1411	0.0026
4	Late (31-120 days)	6900	0.1587	0.0011
3	Default	832	0.1743	0.0034
1	Charged Off	42105	0.1833	0.0005

I flagged all statuses above Late (16-30 days), or $DTIR \lesssim 0.15$, as good/non-defaulted ($good = 1$) and those below, or $DTIR \gtrsim 0.15$, as bad/defaulted ($good = 0$) loans.

This results in 416 002 $good = 1$ loans and 49837 $good = 0$, which differ slightly from the UdeMY numbers (see Sect. 4.4), which are from an unspecified regression model.

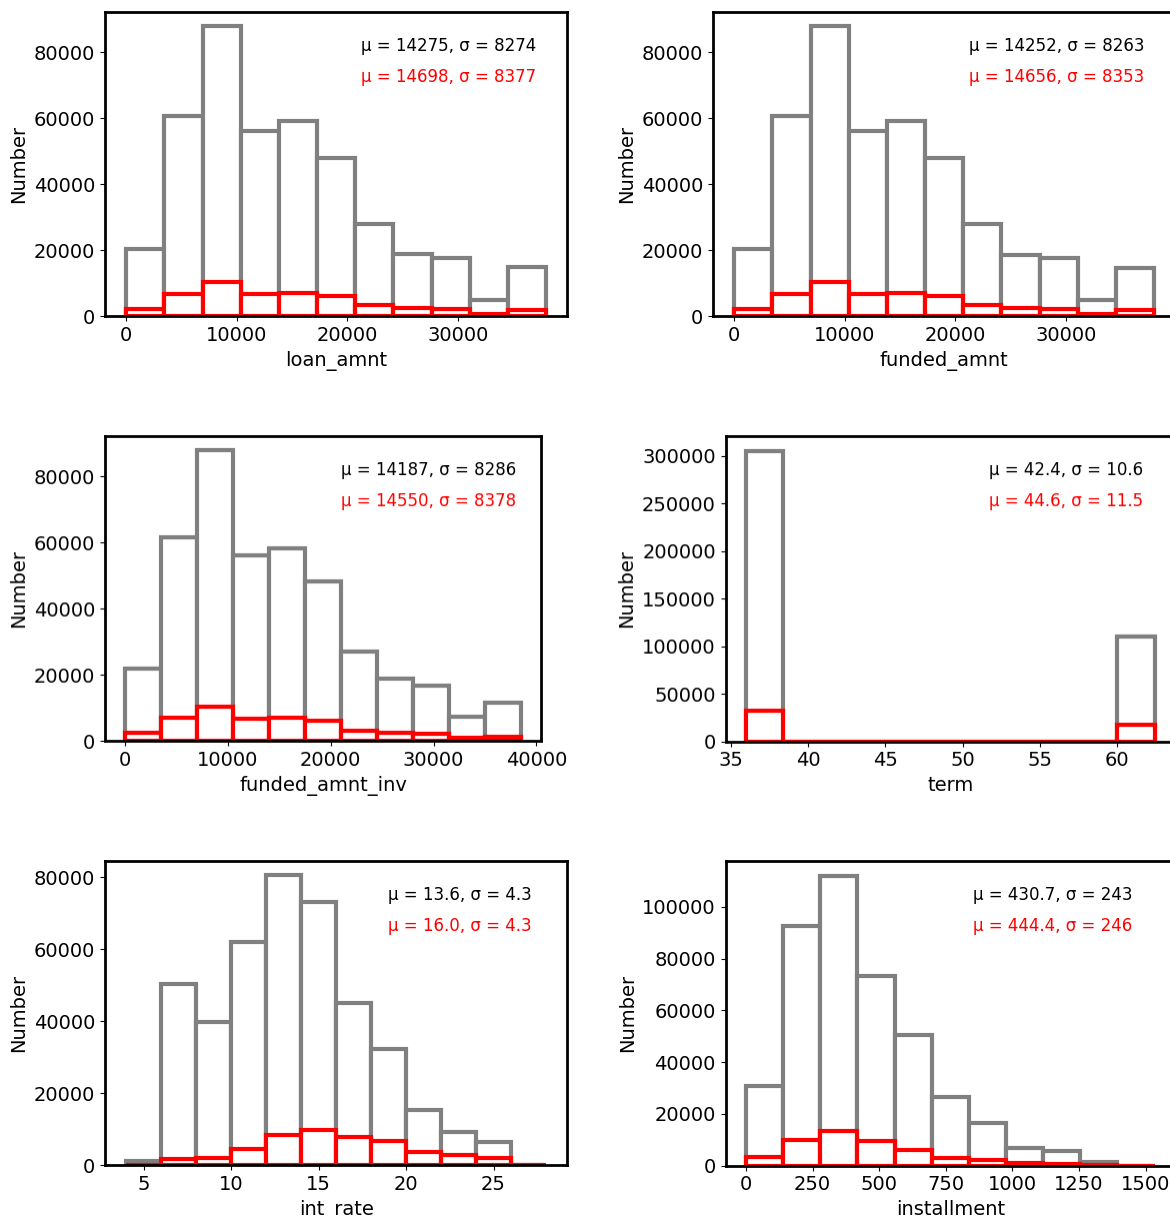
	<code>loan_status</code>	<code>good</code>
0	Fully Paid	1
1	Charged Off	0
2	Fully Paid	1
...
465836	Current	1
465837	Fully Paid	1
465838	Current	1

Since this completes the first phase of the pre-processing and the code is becoming cumbersome, the currently processed data are saved to `loan_data_2007_2014_1.csv`.


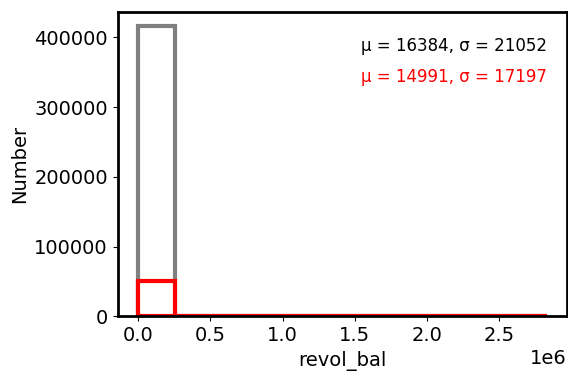

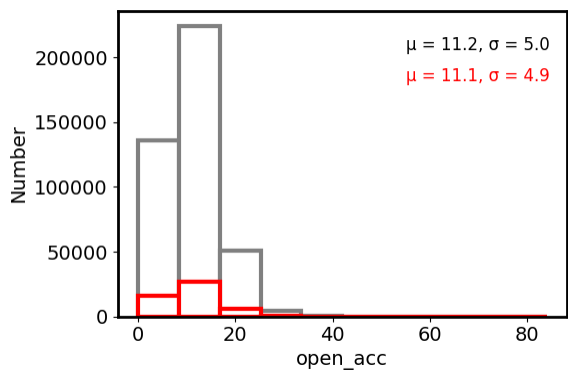

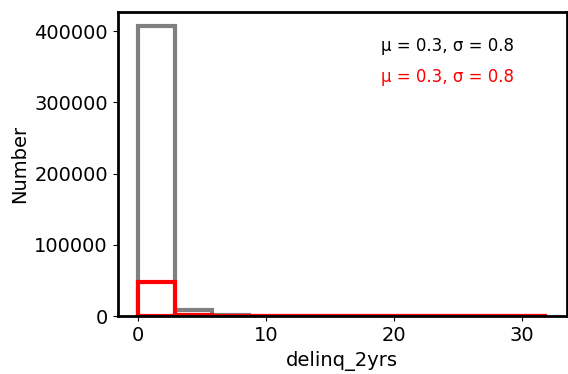

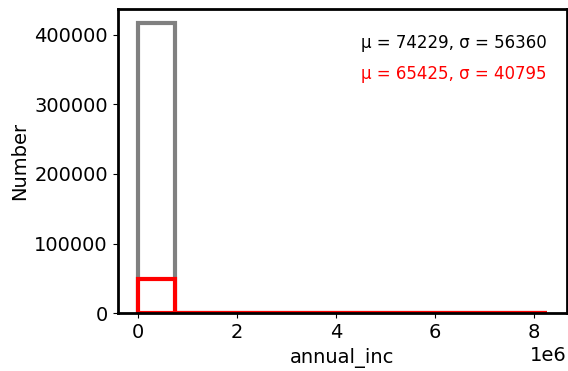
This file is read into `feature_histos.py` to show the distributions for each of the numeric features according to good (black) and bad (red histogram) loans. From these, the good and bad loans mostly overlap, although









³Note that the feature `dti` is not the debt-to-income ratio, but, according to [Kaggle](#), a ratio of the borrower's total monthly debt payments on the total debt obligation.



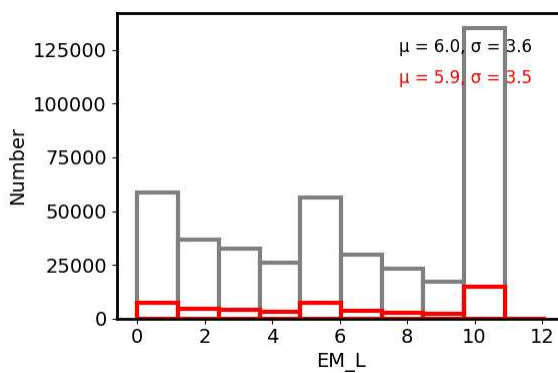
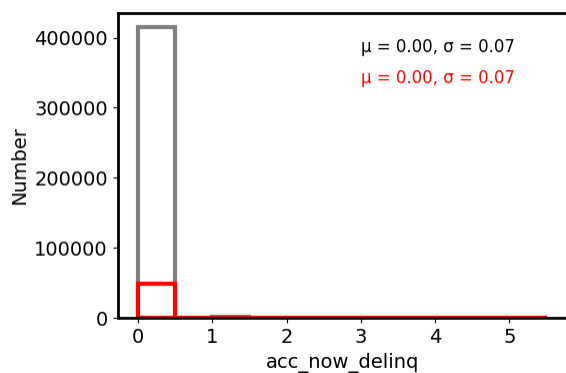
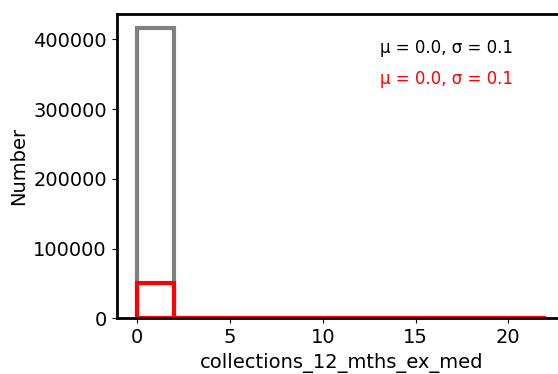



recoveries, collection_recovery_fee, LP_m, debt and, not surprisingly⁴, DTIR appear distinctive, the standard deviations in the data are very wide. This demonstrates that bad loans would be difficult to distinguish based upon the distributions alone.

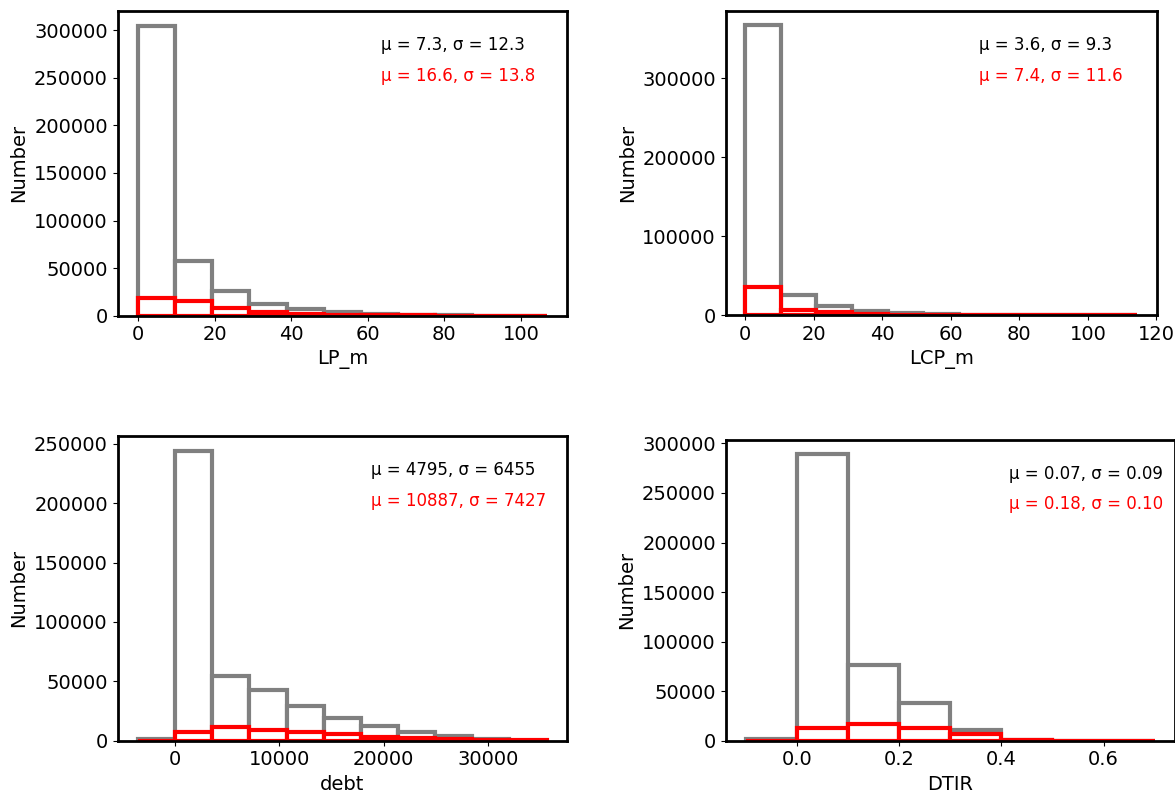


⁴Since this was used to classify the loan as being good or bad.









2.2 Coarse classing of categorical values

Although, unlike the standard practice (e.g. [Udemy](#)), I will not convert the continuous variables to dummy variables, this should be done for the categorical values in order to fully utilise the available data. I will then assess the feature's explanatory power (information value) with respect to the borrower classification.

The remaining categorical fields are

```
['grade', 'home_ownership', 'verification_status', 'pymnt_plan', 'purpose',
 'addr_state', 'initial_list_status'],
```

which are worked through one-by-one in `pre-process_2.py`, below.

2.2.1 Grade

The grade feature, which represents an external grade showing the credit worthiness of the borrower, ranging from A (the highest) to G (the lowest), has only seven unique values.

Here and for the following categories, as per the standard practice, we wish to know how relevant the feature

is to the target variable. For this the *weight of evidence* for each value of the feature is used,

$$WoE_i \equiv \ln \left(\frac{\%(y=1)_i}{\%(y=0)_i} \right) = \ln \left(\frac{\% \text{ good}_i}{\% \text{ bad}_i} \right),$$

which is the natural logarithm of the ratio of proportion of observations in the first target category to those in the second.

Multiplying the *WoE* by the difference in the proportions and summing, gives the *information value*

$$IV = \sum_{i=1}^k \left[(\% \text{ good} - \% \text{ bad}) \times \ln \left(\frac{\% \text{ good}}{\% \text{ bad}} \right) \right],$$

which quantifies how much information the feature contributes to explaining the target. The information value ranges from $IV = 0$ to 1 according to its predictive power (summarised in the table).

Information value	Predictive power
$IV < 0.02$	None
$0.02 < IV < 0.1$	Weak
$0.1 < IV < 0.3$	Medium
$0.3 < IV < 0.5$	Strong
$IV > 0.5$	Suspect

Proceeding with the grade feature, in increasing *WoE*

```
n_obs - number of observations for this grade
%n_obs - percentage of total number of observations
n_good - number of good = 1 observations for this grade (bad loans are flagged 'good = 0')
%good - percentage of good = 1 observations for this grade
%n_good - percentage of total number
```

	grade	n_obs	%n_obs	n_good	n_bad	%good	%bad	%n_good	%n_bad	WoE	IV
0	G	3316	0.71	2470	846	74.49	25.51	0.59	1.70	-1.06	0.011766
1	F	13205	2.83	10129	3076	76.71	23.29	2.43	6.17	-0.93	0.034782
2	E	35679	7.66	28945	6734	81.13	18.87	6.96	13.51	-0.66	0.043230
3	D	76784	16.48	65228	11556	84.95	15.05	15.68	23.19	-0.39	0.029289
4	C	125186	26.87	111042	14144	88.70	11.30	26.69	28.38	-0.06	0.001014
5	B	136848	29.38	126233	10615	92.24	7.76	30.34	21.30	0.35	0.031640
6	A	74821	16.06	71955	2866	96.17	3.83	17.30	5.75	1.10	0.127050

```
=====
For 'grade' feature, sum of IVs = 0.2788: 0.1 < IV < 0.3 - medium predictive power
=====
```

The predictive power of the grade feature is close to the medium/strong threshold.

Since there are only seven different values, each with a large number of observations, the dummy variables can be one-hot-encoded directly, thus not requiring coarse classing, giving, for example,

	good	grade_A	grade_B	grade_C	grade_D	grade_E	grade_F	grade_G
0	1	0	1	0	0	0	0	0
1	0	0	0	1	0	0	0	0

```

2          1          0          0          1          0          0          0          0
...      ...      ...      ...      ...      ...      ...      ...      ...
465837    1          1          0          0          0          0          0          0
465838    1          0          0          0          1          0          0          0

```

2.2.2 Home ownership


Since the *WoE* and *IV* calculations are in a function (in `pre-process_2.py`), it is trivial to get these metrics for the remaining features. Continuing with `home_ownership`

	home_ownership	n_obs	%n_obs	n_good	n_bad	%good	%bad	%n_good	%n_bad	WoE	IV
0	OTHER	182	0.04	155	27	85.16	14.84	0.04	0.05	-0.22	0.000022
1	RENT	188208	40.40	165019	23189	87.68	12.32	39.67	46.53	-0.16	0.010976
2	NONE	46	0.01	39	7	84.78	15.22	0.01	0.01	0.00	0.000000
3	OWN	41654	8.94	37231	4423	89.38	10.62	8.95	8.87	0.01	0.000008
4	MORTGAGE	235748	50.61	213557	22191	90.59	9.41	51.34	44.53	0.14	0.009534
5	ANY	1	0.00	1	0	100.00	0.00	0.00	0.00	NaN	NaN

=====

For 'home_ownership' feature, sum of IVs = 0.0205: 0.02 < IV < 0.1 - weak predictive power

=====



If using this feature, given the low *IV*, the small numbers (`n_obs`) for OTHER, NONE and ANY means that it is not sensible to one-hot-encode these. Therefore, OTHER is merged with the nearest populous value with a similar *WoE*, RENT. Likewise, NONE and ANY is merged with OWN.

After merging (and renaming for neatness), we are left with the fields `HO_MORTGAGE` `HO_RENT+` `HO_OWN+`, where the '+' designates the inclusion of other groups.

2.2.3 Verification status

There is no information regarding the `verification_status` on either Udemy's [Credit Risk Modelling in Python](#) nor [Kaggle](#).

Running the `woe` function in `pre-process_2.py`, shows that there are only three unique values, each of which has a large, and similar, number of observations.

```

      VS  n_obs  %n_obs  n_good  n_bad  %good  %bad  %n_good  %n_bad  WoE      IV
0      Verified 167920  36.05 146802 21118  87.42 12.58   35.29  42.37 -0.18  0.012744
1 Source Verified 149862  32.17 134444 15418  89.71 10.29   32.32  30.94  0.04  0.000552
2   Not Verified 148057  31.78 134756 13301  91.02  8.98   32.39  26.69  0.19  0.010830
=====
For 'VS' feature, sum of IVs = 0.0241: 0.02 < IV < 0.1 - weak predictive power
=====

```

Therefore, there is no need to coarse class this feature, which are therefore converted to dummy variables directly.

2.2.4 Payment plan

This flags whether a payment plan has been put in place for the loan.

```

PP  n_obs  %n_obs  n_good  n_bad  %good  %bad  %n_good  %n_bad  WoE  IV
0  y      9      0.0      5      4  55.56  44.44      0.0   0.01 -inf  inf
1  n 465830 100.0 415997 49833  89.30 10.70     100.0  99.99  0.0  0.0
=====
For 'PP' feature, sum of IVs = inf: IV > 0.5 - something fishy going on
=====

```

There are just two values, with the number of PP = y being so small that the feature is almost exclusively PP = n. This highly unbalanced feature is therefore removed.

2.2.5 Purpose

This is the purpose of the loan.

```

      P      n_obs  %n_obs  n_good  n_bad  %good  %bad  %n_good  %n_bad  WoE      IV
small_business  6993  1.50   5572  1421  79.68  20.32   1.34  2.85 -0.75  0.011325
moving         2990  0.64   2551  439   85.32  14.68   0.61  0.88 -0.37  0.000999
renewable_energy  349  0.07    298  51   85.39  14.61   0.07  0.10 -0.36  0.000108
house          2264  0.49   1962  302   86.66  13.34   0.47  0.61 -0.26  0.000364
other          23607  5.07  20507  3100  86.87  13.13   4.93  6.22 -0.23  0.002967
educational     419  0.09    366  53   87.35  12.65   0.09  0.11 -0.20  0.000040
medical         4588  0.98   4007  581   87.34  12.66   0.96  1.17 -0.20  0.000420
vacation        2484  0.53   2200  284   88.57  11.43   0.53  0.57 -0.07  0.000028
debt_consolidation 273984 58.82 243389 30595  88.83  11.17  58.51 61.39 -0.05  0.001440
wedding         2331  0.50   2071  260   88.85  11.15   0.50  0.52 -0.04  0.000008
home_improvement 26510  5.69  24001  2509  90.54  9.46   5.77  5.03  0.14  0.001036
major_purchase  9813  2.11   8911  902   90.81  9.19   2.14  1.81  0.17  0.000561
car             5389  1.16   4929  460   91.46  8.54   1.18  0.92  0.25  0.000650


```



```

credit_card 104118 22.35 95238 8880 91.47 8.53 22.89 17.82 0.25 0.012675
=====
For 'P' feature, sum of IVs = 0.0326: 0.02 < IV < 0.1 - weak predictive power
=====

```



There are some low numbers of observations (e.g. `renewable_energy`), which could be merged with the other values (as in Sect. 2.2.2), but given that there are 14 unique values, these should be bundled into coarser classes in any case.

With the guidance of the *WoE* values in the above table and the plot, the dummies are combined as

```

comb_dummy('P', ['small_business', 'moving', 'renewable_energy'])
comb_dummy('P', ['other', 'house', 'educational', 'medical'])
comb_dummy('P', ['debt_consolidation', 'vacation', 'wedding'])
comb_dummy('P', ['home_improvement', 'major_purchase'])
comb_dummy('P', ['credit_card', 'car']);

```

leaving the features

```

P_small_business+ P_other+ P_debt_consolidation+ P_home_improvement+ P_credit_card+

```

2.2.6 Address state

	AS	n_obs	%n_obs	n_good	n_bad	%good	%bad	%n_good	%n_bad	WoE	IV
0	NE	14	0.00	9	5	64.29	35.71	0.00	0.01	-inf	inf
1	NV	6508	1.40	5641	867	86.68	13.32	1.36	1.74	-0.25	0.000950
2	HI	2485	0.53	2175	310	87.53	12.47	0.52	0.62	-0.18	0.000180

3	FL	31609	6.79	27738	3871	87.75	12.25	6.67	7.77	-0.15	0.001650
4	AL	5849	1.26	5138	711	87.84	12.16	1.24	1.43	-0.14	0.000266
5	NY	40189	8.63	35585	4604	88.54	11.46	8.55	9.24	-0.08	0.000552
6	LA	5483	1.18	4853	630	88.51	11.49	1.17	1.26	-0.07	0.000063
7	OK	4111	0.88	3644	467	88.64	11.36	0.88	0.94	-0.07	0.000042
8	NM	2586	0.56	2292	294	88.63	11.37	0.55	0.59	-0.07	0.000028
9	NC	12669	2.72	11233	1436	88.67	11.33	2.70	2.88	-0.06	0.000108
10	NJ	18039	3.87	16018	2021	88.80	11.20	3.85	4.06	-0.05	0.000105
11	MD	10959	2.35	9735	1224	88.83	11.17	2.34	2.46	-0.05	0.000060
12	VA	14211	3.05	12616	1595	88.78	11.22	3.03	3.20	-0.05	0.000085
13	CA	71369	15.32	63425	7944	88.87	11.13	15.25	15.94	-0.04	0.000276
14	MO	7500	1.61	6669	831	88.92	11.08	1.60	1.67	-0.04	0.000028
15	UT	3428	0.74	3052	376	89.03	10.97	0.73	0.75	-0.03	0.000006
16	TN	5980	1.28	5323	657	89.01	10.99	1.28	1.32	-0.03	0.000012
17	MI	11542	2.48	10283	1259	89.09	10.91	2.47	2.53	-0.02	0.000012
18	AZ	10702	2.30	9534	1168	89.09	10.91	2.29	2.34	-0.02	0.000010
19	AR	3488	0.75	3108	380	89.11	10.89	0.75	0.76	-0.01	0.000001
20	SD	980	0.21	876	104	89.39	10.61	0.21	0.21	0.00	0.000000
21	RI	2049	0.44	1828	221	89.21	10.79	0.44	0.44	0.00	0.000000
22	PA	16416	3.52	14660	1756	89.30	10.70	3.52	3.52	0.00	0.000000
23	OH	15228	3.27	13605	1623	89.34	10.66	3.27	3.26	0.00	0.000000
24	KY	4433	0.95	3967	466	89.49	10.51	0.95	0.94	0.01	0.000001
25	MN	8150	1.75	7292	858	89.47	10.53	1.75	1.72	0.02	0.000006
26	IN	6523	1.40	5841	682	89.54	10.46	1.40	1.37	0.02	0.000006
27	MA	11058	2.37	9913	1145	89.65	10.35	2.38	2.30	0.03	0.000024
28	DE	1270	0.27	1139	131	89.69	10.31	0.27	0.26	0.04	0.000004
29	WA	10508	2.26	9455	1053	89.98	10.02	2.27	2.11	0.07	0.000112
30	GA	14956	3.21	13451	1505	89.94	10.06	3.23	3.02	0.07	0.000147
31	OR	5946	1.28	5357	589	90.09	9.91	1.29	1.18	0.09	0.000099
32	WI	5908	1.27	5329	579	90.20	9.80	1.28	1.16	0.10	0.000120
33	TX	36398	7.81	32999	3399	90.66	9.34	7.93	6.82	0.15	0.001665
34	CT	7196	1.54	6537	659	90.84	9.16	1.57	1.32	0.17	0.000425
35	IL	18602	3.99	16890	1712	90.80	9.20	4.06	3.44	0.17	0.001054
36	SC	5580	1.20	5077	503	90.99	9.01	1.22	1.01	0.19	0.000399
37	CO	9735	2.09	8861	874	91.02	8.98	2.13	1.75	0.20	0.000760
38	AK	1251	0.27	1143	108	91.37	8.63	0.27	0.22	0.20	0.000100
39	MT	1394	0.30	1271	123	91.18	8.82	0.31	0.25	0.22	0.000132
40	KS	4186	0.90	3817	369	91.18	8.82	0.92	0.74	0.22	0.000396
41	NH	2230	0.48	2048	182	91.84	8.16	0.49	0.37	0.28	0.000336
42	VT	904	0.19	827	77	91.48	8.52	0.20	0.15	0.29	0.000145
43	MS	1224	0.26	1124	100	91.83	8.17	0.27	0.20	0.30	0.000210
44	WV	2410	0.52	2220	190	92.12	7.88	0.53	0.38	0.33	0.000495
45	WY	1128	0.24	1045	83	92.64	7.36	0.25	0.17	0.39	0.000312
46	DC	1425	0.31	1331	94	93.40	6.60	0.32	0.19	0.52	0.000676
47	IA	14	0.00	13	1	92.86	7.14	0.00	0.00	NaN	NaN

```


48 ID      12      0.00      11      1  91.67  8.33      0.00      0.00  NaN      NaN
49 ME       4      0.00       4      0 100.00  0.00      0.00      0.00  NaN      NaN

```

```

=====
For 'AS' feature, sum of IVs = 0.0121: IV < 0.02 - no predictive power
=====

```



Again, there are values which have a very low number of observations (e.g. NE), as well 50 different classes. These are therefore merged, according to `n_obs` and `WoE`, as

```

comb_dummy('AS', ['FL', 'NV', 'HI', 'NE', 'AL'])      AS_CA 71369      AS_WA 10508
comb_dummy('AS', ['LA', 'OK', 'NM'])                  AS_MD 10959      AS_FL+ 46465
comb_dummy('AS', ['MO', 'UT', 'TN'])                  AS_MI 11542      AS_LA+ 12180
comb_dummy('AS', ['AZ', 'AR', 'SD', 'RI'])             AS_NC 12669      AS_MO+ 16908
comb_dummy('AS', ['KY', 'MN', 'IN', 'MA', 'DE'])      AS_NJ 18039      AS_KY+ 31434
comb_dummy('AS', ['GA', 'OR', 'WI'])                  AS_NY 40189      AS_GA+ 26810
comb_dummy('AS', ['CT', 'IL'])                        AS_OH 15228      AS_CT+ 25798
comb_dummy('AS', ['SC', 'CO', 'AK', 'MT', 'KS'])      AS_PA 16416      AS_SC+ 22146
comb_dummy('AS', ['NH', 'VT', 'MS', 'WV'])            AS_TX 36398      AS_NH+ 6768
comb_dummy('AS', ['WY', 'DC', 'IA', 'ID', 'ME']);     AS_VA 14211      AS_WY+ 2583

```

where the resulting classes and `n_obs` are shown in the two right-hand columns. Arguably `AS_NH+` and `AS_WY+` could also be merged to increase the size of this class, but with $WoE = 0.28 - 0.33$, cf. ≥ 0.39 , respectively, this is left as is.

2.2.7 Initial list status

This is the initial listing status of the loan.

	ILS	n_obs	%n_obs	n_good	n_bad	%good	%bad	%n_good	%n_bad	WoE	IV
0	f	302666	64.97	267251	35415	88.30	11.70	64.24	71.06	-0.10	0.006820
1	w	163173	35.03	148751	14422	91.16	8.84	35.76	28.94	0.21	0.014322

=====

For 'ILS' feature, sum of IVs = 0.0211: 0.02 < IV < 0.1 - weak predictive power

=====

With only two values, each with a large number of observations, this can be converted directly to dummy variables.

Now that the second phase of the pre-processing is complete, the fully processed data are saved (to `loan_data_2007_2014_2.csv`).

3 Machine Learning

3.1 Initial models and tests

The standard practice is to model the *probability of default* (PD) with the logistic regression classifier. Since, it is straightforward to run other common, but different, binary classifiers in tandem with this, I include these in the modelling. That is, the data are tested with:

1. *Logistic Regression* (LR)
2. *k-Nearest Neighbour* (kNN)
3. *Support Vector Classifier* (SVC)
4. *Decision Tree Classifier* (DTC)


Each of which is described in further detail by [Curran \(2021\)](#).

The machine learning is performed in `ML.py`, where the imbalance in the data has to be addressed. Specifically, there are 49 782 bad loans compared to $465\,839 - 49\,837 = 416\,057$ good loans, and so from raw statistics alone we can ascertain a $100 \times 416\,057 / 465\,839 = 89.3\%$ probability of the loan being good, although this would have no predictive power. Therefore 49 837 of the good loan from the pool of 416 057 are randomly selected for the binary testing.

Training on 80% and testing on 20% of the sample, using (close to) the default standard parameters (`n_neighbors = 10` and `max_depth = 5`), high scores are obtained, especially for the decision tree classifier.

For a 0.2 test fraction (79651 train & 19913 test) LR training score = 92.418
 For a 0.2 test fraction (79651 train & 19913 test) KNN training score = 80.765
 For a 0.2 test fraction (79651 train & 19913 test) SVC training score = 93.851
 For a 0.2 test fraction (79651 train & 19913 test) DTC training score = 98.463

It should be noted, however, that this is for a single run and that the scores to exhibit some variation. Also, high scores may not be unexpected since the target was defined via the `loan_status`, i.e. `good = 1` where $DTIR < 0.15$ and `good = 0` where $DTIR > 0.15$ (Sect. 2.1.5).





Using `target_classes.py`, the above plot shows the debt-to-income ratio distributions for each status. These overlap to such an extent that any prediction of the status based upon the DTIR alone would essentially be guesswork.

3.2 Model optimisation

Although the SVC scores relatively high, it is very CPU intensive for this size of dataset and so we proceed with the other three classifiers.

Using `LR.py`, for the LR classifier a grid search was also CPU intensive and failed to complete. Manual experimentation honed in on `C=10`, `solver='lbfgs'` as close to optimal.



Using `kNN.py` and `DTC.py` for the kNN and DT classifiers, the number of nearest neighbours and maximum depths were iterated, giving `n_neighbors` ≈ 2300 (a validation score of 86.4%) and `max_depth` = 10


(98.5%), respectively.

The confusion matrix for a run of the DT classifier is shown to the right. From the validation of the model, TP is the number of true positives, FP – false positives, FN – false negatives and TN – true negatives.

	Predicted			
	Good	Bad		
Actual good	9734	90	TP	FP
Actual bad	225	9864	FN	TN

The validation data yield 9734 true positives and 9864 true negatives, compared to 90 false positives and 225 false negatives, thus giving a validation score of $(9734 + 9864)/(9734 + 9864 + 90 + 225) = 0.984$.

This is an exceedingly good score, although it was noted that different runs could yield quite different results upon randomisation of the data. In order to check that the scores are representative we ran each classifier 100 times randomising the sample of good = 1 selected from the pool of 416 057 good loans each time.



From the resulting distribution of accuracies, it is seen that, although the DT classifier can score very high, its accuracy shows a strong dependency on the data being used, suggestive of over-fitting. On the other hand, the LR classifier remains stable with a mean accuracy of $92.4 \pm 0.2\%$.

3.3 Feature importance

Proceeding with the LR classifier, in decreasing positive importance the most important features are:

	Feature	Importance
18	total_pymnt_inv	2.596452e-01
2	funded_amnt_inv	2.097149e-01
16	out_prncp_inv	1.924370e-01
19	total_rec_prncp	1.428808e-01
15	out_prncp	1.159634e-01
20	total_rec_int	7.853511e-02
22	recoveries	7.596138e-02
0	loan_amnt	7.492436e-02
5	installment	5.050784e-02
31	LP_m	4.075278e-02
33	debt	2.631731e-02
17	total_pymnt	2.428846e-02
24	last_pymnt_amnt	1.451834e-02
1	funded_amnt	1.387553e-02
21	total_rec_late_fee	1.193456e-02

Eliminating the least important features (in `ML_feat.py`), we see that the validation accuracy only drops below 90% when just the top three features (`total_pymnt_inv`, `funded_amnt_inv` & `out_prncp_inv`) are used, all of which are numerical.

No.	Last feature used	Feature importance	Score [%]
15	total_rec_late_fee	> 0.01	92.6
9	installment	> 0.05	90.8
5	out_prncp	> 0.1	91.2
3	out_prncp_inv	$\gtrsim 0.2$	87.9

Given that the prediction accuracy using only the 15 most important features is as high as when using all the features, much of the above processing, including the coarse classing, is unnecessary (at least in this case). This may be evident via the low information values for most of the categorical variables (Sect. 2.2.)


4 Results

4.1 ROC curve

The *receiver operating characteristic* (ROC) is a plot of the *true positive rate* versus the *false positive rate*, obtained from the confusion matrix at a given classification threshold (e.g. $> 0.5 \Rightarrow \text{good} = 1$, $< 0.5 \Rightarrow \text{good} = 0$).

The area under area under the ROC curve (AUC) provides a measure of the effectiveness of the classification. The scores are grouped into categories in the table to the right. For example, predicting purely by chance gives the $\text{AUC} = 0.5$, shown by the red broken lines in the following plot.

Area under ROC curve	Model Interpretation
0.5 – 0.6	No discrimination
0.6 – 0.7	Poor classifier
0.7 – 0.8	Fair classifier
0.8 – 0.9	Good classifier
0.9 – 1	Excellent classifier



Using all of the features, the LR classifier gives an $AUC = 0.974$, scoring much higher the 0.5 expected from chance and significantly higher than Udemy's [Credit Risk Modelling in Python](#) AOC of 0.702.

According to [Credit Scoring Series Part Five](#), "In credit risk, an AUC of 0.75 or higher is the industry-accepted standard and prerequisite to model acceptance".

4.2 Gini coefficient

The Gini coefficient measures the inequality between the good = 1 and good = 0 borrowers. This is obtained from the AUC, via


$$\text{Gini} = (2 \times \text{AUC}) - 1,$$

so that the higher the Gini coefficient the better the model in distinguishing the two classes.

The above AUC gives $\text{Gini} = 94.7\%$, compared to 40%, from the same data, by [Udemy](#) and 71% by [Credit Risk Modelling \(Part II\)](#).

4.3 Kolmogorov-Smirnov score

The Kolmogorov-Smirnov (KS) test compares the distribution of a borrower's credit score with the reference distribution. It uses the maximum difference between the cumulative distribution functions to determine whether the two distributions are significantly different from each other.



In the plot the blue trace represents the good = 1 borrowers and the red trace the good = 0 borrowers. The score of 0.85 is high, indicating that the model has excellent predictive power. Note that both [Udemy](#) and [Credit Risk Modelling \(Part II\)](#) obtain a score of just 0.30.

4.4 Using the Udemmy Target Classification

Given the results are much superior to those of [Udemmy](#), in order to check that these are not due to the definition of the target (Sect. 2.1.5), we repeat the analysis using the Udemmy classification.

Proceeding with [pre-process_Udemmy_1.py](#), the two classification schemes are shown below.

<pre> ===== my_good = 1 ===== ['Fully Paid' 'Current' 'In Grace Period' 'Late (16-30 days)' 'Does not meet the credit policy. Status:Fully Paid' 'Does not meet the credit policy. Status:Charged Off'] ===== my_good = 0 ===== ['Charged Off' 'Default' 'Late (31-120 days)'] Giving 416002 good loans (good = 1) and 49837 bad loans (good = 0) mean 0.893017 std 0.309092 min 0.000000 max 1.000000 </pre>	<pre> ===== Udemmy good = 1 ===== ['Fully Paid' 'Current' 'In Grace Period' 'Late (16-30 days)' 'Does not meet the credit policy. Status:Fully Paid'] ===== Udemmy good = 0 ===== ['Charged Off' 'Default' 'Late (31-120 days)' 'Does not meet the credit policy. Status:Charged Off'] Giving 415256 good loans (good = 1) and 50583 bad loans (good = 0) mean 0.891415 std 0.311118 min 0.000000 max 1.000000 </pre>
--	--

These are similar, with only `Does not meet the credit policy. Status:Charged Off` being classed differently.

The previous steps (Sects. 2.1.2 – 2.2.7) are repeated for the Udemmy classification, with the coarse classing done in [pre-process_Udemmy_2.py](#) and the machine learning in [ML_Udemmy.py](#), where a randomly selected sample of 50583 good = 1 loans is used to match the number of good = 0 loans.


A typical run of the various classifiers gives

For a 0.2 test fraction (80932 train & 20234 test) LR training score = 92.947 percent
Validation accuracy of LR is 92.83 percent

For a 0.2 test fraction (80932 train & 20234 test) KNN training score = 84.062 percent
Validation accuracy of KNN is 84.55 percent

For a 0.2 test fraction (80932 train & 20234 test) DTC training score = 98.575 percent
Validation accuracy of DTC is 78.28 percent

which are very similar scores to those above. Therefore, the slightly different classification of loan status is not responsible for the higher scores.



The AUC, Gini coefficient (95.39%) and KS score are also very similar using this classification.

4.5 IVs of the Most Important Features

It has been shown above that retaining the continuous features results in a much superior predictive model, with a Gini coefficient of 95% (cf. 40% with fine classing) and a KS score of 0.85 (cf. 0.30).

Furthermore, validation scores of $\geq 90\%$ are retained just using the five most important features. To see how these compare with the categorical features, which were coarse classed, the weight of evidence and information value is examined.


The most important feature is `total_pymnt_inv`, which [Kaggle](#) defines as “*payments received to date for portion of total amount funded by investors*”. This feature ranges from 0 to 57 777.58 with 34 7341 discrete values. This is therefore binned into intervals of 5000, giving the *WoEs*.

```
===== total_pymnt_inv (TPI) =====
```

	TPI_bin	n_obs	%n_obs	n_good	n_bad	%good	%bad	%n_good	%n_bad	WoE	IV
0	5000	153440	35.07	130617	22823	85.13	14.87	32.88	56.83	-0.55	0.131725
1	10000	123710	28.28	113400	10310	91.67	8.33	28.54	25.67	0.11	0.003157
2	15000	72279	16.52	67972	4307	94.04	5.96	17.11	10.72	0.47	0.030033
3	20000	39705	9.08	38070	1635	95.88	4.12	9.58	4.07	0.86	0.047386
4	25000	22822	5.22	22152	670	97.06	2.94	5.58	1.67	1.21	0.047311
5	30000	12029	2.75	11786	243	97.98	2.02	2.97	0.61	1.58	0.037288
6	35000	6461	1.48	6345	116	98.20	1.80	1.60	0.29	1.71	0.022401
7	40000	4815	1.10	4771	44	99.09	0.91	1.20	0.11	2.39	0.026051
8	45000	1753	0.40	1738	15	99.14	0.86	0.44	0.04	2.40	0.009600
9	50000	394	0.09	394	0	100.00	0.00	0.10	0.00	inf	inf
10	55000	59	0.01	59	0	100.00	0.00	0.01	0.00	inf	inf


```
=====
```

For 'TPI_bin' feature, sum of IVs = 0.3550: 0.3 < IV < 0.5 - strong predictive power



The next most important feature is `funded_amnt_inv`, which [Kaggle](#) defines as “the total amount committed by investors for that loan at that point in time”. This has 9854 unique values ranging from 0 to 35 000, so again the data are binned in intervals of 5000.

```
===== funded_amnt_inv (FAI) =====
FAI_bin  n_obs  %n_obs  n_good  n_bad  %good  %bad  %n_good  %n_bad  WoE      IV
4   30000  23269   5.36   20627  2642   88.65  11.35    5.32    5.69 -0.07  0.000259
2   20000  67354  15.52   59799  7555   88.78  11.22   15.43   16.27 -0.05  0.000420
3   25000  36385   8.38   32394  3991   89.03  10.97    8.36    8.60 -0.03  0.000072
1   15000  81262  18.72   72444  8818   89.15  10.85   18.69   18.99 -0.02  0.000060
0   10000 132276  30.47  118343 13933   89.47  10.53   30.53   30.01  0.02  0.000104
5    5000  93531  21.55   84039  9492   89.85  10.15   21.68   20.44  0.06  0.000744
=====
For 'FAI_bin' feature, sum of IVs = 0.0017: IV < 0.02 - no predictive power
=====
```



The third feature is `out_prncp_inv`, which [Kaggle](#) defines as “the remaining outstanding principal for portion of total amount funded by investors”. This has 141 181 unique values ranging from 0 to 32 160.38, so again the data are binned in intervals of 5000.


```
===== out_prncp_inv (OPI) =====
```

	OPI_bin	n_obs	%n_obs	n_good	n_bad	%good	%bad	%n_good	%n_bad	WoE	IV
4	30000	2638	1.34	2463	175	93.37	6.63	1.29	2.52	-0.67	0.008241
3	25000	7656	3.88	7307	349	95.44	4.56	3.84	5.03	-0.27	0.003213
2	20000	16349	8.29	15670	679	95.85	4.15	8.23	9.79	-0.17	0.002652
1	15000	31703	16.07	30514	1189	96.25	3.75	16.03	17.14	-0.07	0.000777
0	10000	56365	28.57	54405	1960	96.52	3.48	28.58	28.25	0.01	0.000033
5	5000	82600	41.86	80013	2587	96.87	3.13	42.03	37.28	0.12	0.005700

```
=====
```

```
For 'OPI_bin' feature, sum of IVs = 0.0206: 0.02 < IV < 0.1 - weak predictive power
```

```
=====
```



Interestingly, only `total_pymnt_inv` has any real predictive power, but if it was as straightforward as building an accurate predictive model based upon information values alone, there would be no need for machine learning.